

# **American Express - Default Prediction**

## **MIDS - W207**

### **Spring 2023**

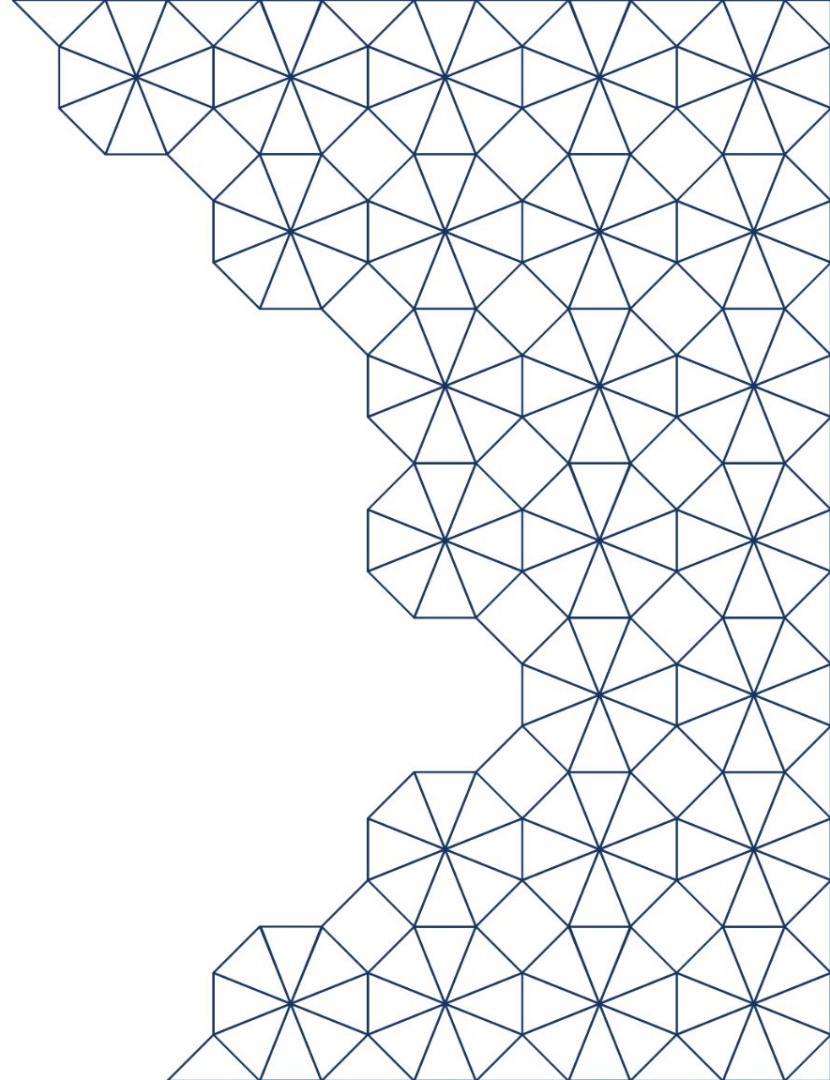
Jim Zhu

Rick Chen

Julian Rippert

# Agenda

1. Problem Motivation
2. Dataset Description
3. EDA
4. Models
  - a. Clustering
  - b. Random Forest
  - c. Transformer
5. Results
6. Takeaways



## **Problem Motivation**

- Goal of a credit issuer is to ensure credit is repaid back with interest to compensate risk
- Default events are unfavorable for a lender

The American Express logo is displayed in a bold, blue, sans-serif font. The word "AMERICAN" is on the top line and "EXPRESS" is on the bottom line, both in all caps.

# Dataset Description

- Dataset dimensions
  - (5531451,191)
- Time Series Data
- Default is when a customer does not pay due amount in 120 days after their latest statement date.
- *Features are anonymized and normalized*

## Feature Categories:

$D_*$  = Delinquency variables

$S_*$  = Spend variables

$P_*$  = Payment variables

$B_*$  = Balance variables

$R_*$  = Risk variables

# EDA

- Validated the data was standardized
- Categoricals Columns
  - Reindexed to 0
  - One-Hot Encoded
  - Dropped due to interpretability

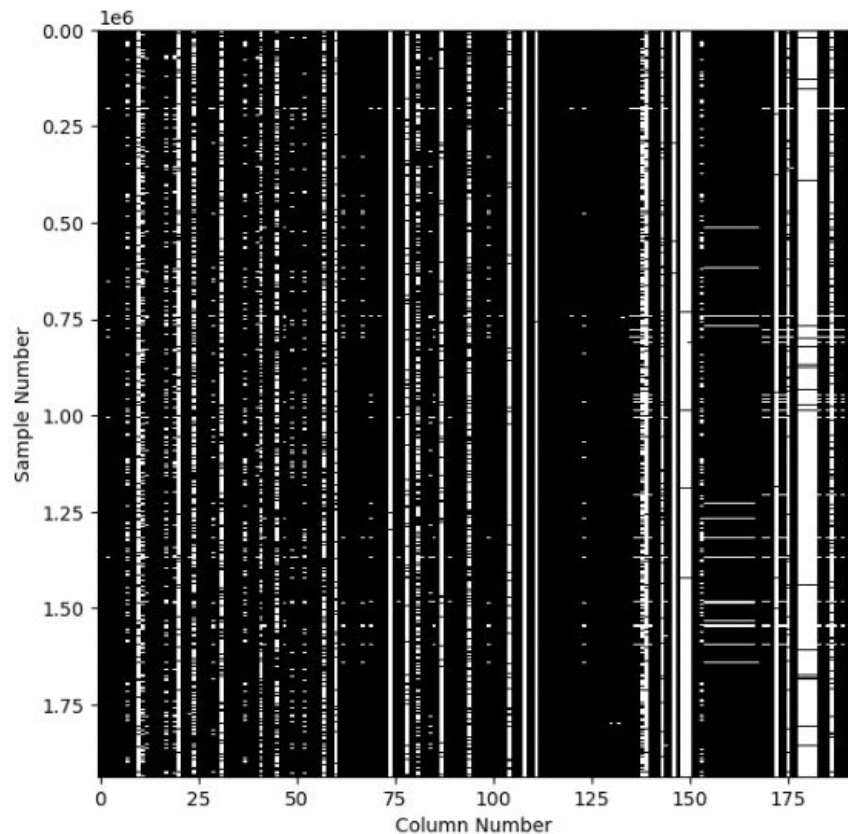
	count	mean	std	min	25%	50%	75%	max
D_42	283,952.000000	0.000000	0.000000	-0.000372	0.040314	0.123779	0.258057	4.187500
D_49	171,995.000000	0.000000	0.000000	0.000007	0.062439	0.131104	0.245972	19.953125
D_53	477,308.000000	0.000000	0.000000	0.000000	0.005875	0.011810	0.039185	7.902344
R_4	1,936,007.000000	0.000000	0.000000	0.000000	0.002539	0.005081	0.007626	1.009766
R_5	1,936,007.000000	0.000000	0.000000	0.000000	0.002550	0.005096	0.007645	17.515625

```
2.0    726993
3.0    442452
1.0    390498
5.0    147417
4.0     94695
7.0     81702
6.0     51298
Name: B_38, dtype: int64
```

```
CO     1428972
CR      336769
CL     154494
XZ       9405
XM       3672
XL       2695
Name: D_63, dtype: int64
```

```
O     1022626
U      514145
R      294146
      69495
-1      35595
Name: D_64, dtype: int64
```

# EDA - NaNs



- Set max 40% threshold of NaNs per feature
  - Length before filtering out high NaN count cols: 191
  - Number of columns to drop: 62
- Imputed -1 for all NaNs

# Train/Val/Test Split

```
Training Features dimensions: (3318870, 190)
Training Labels dimensions: (3318870,)
Validation Features dimensions: (1106290, 190)
Validation Labels dimensions: (1106290,)
Test Features dimensions: (1106291, 190)
Test Labels dimensions: (1106291,)
```

```
Training Percent of Positive Labels: 0.2440
```

```
Validation Percent of Positive Labels: 0.2553
```

- Implemented a 60/20/20 split
- To maintain integrity of the time series no shuffling was used
- Yet, similar target label distribution

## Models & Experiments

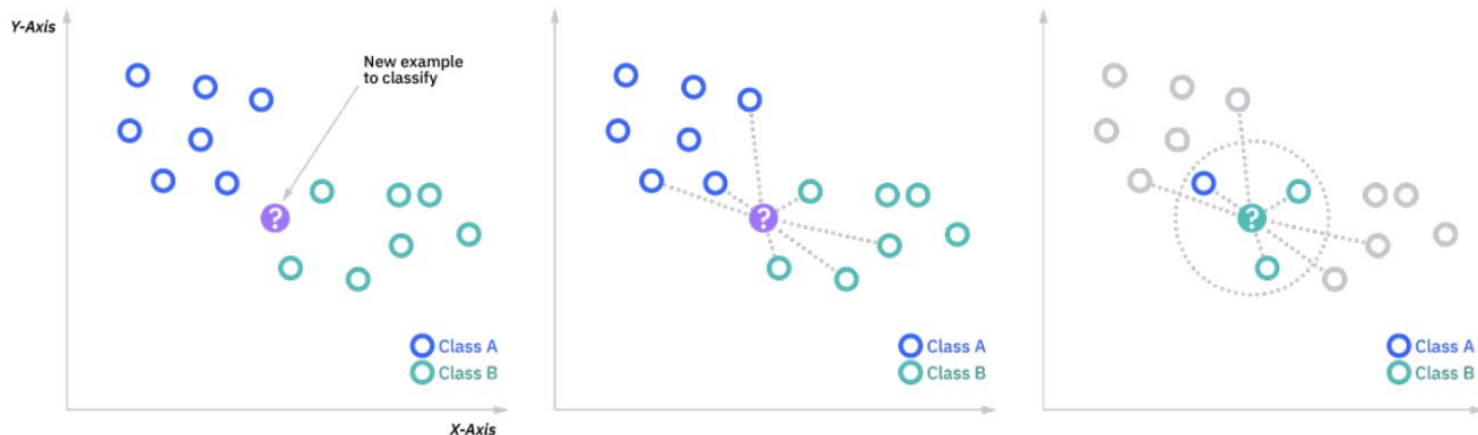
- Baseline Model - always predict no default event
  - Accuracy Score = 75.60%



# Model - KNN

KNN models are relatively simple models great for both classification and regression problems

- Known as a lazy learner, and makes a great baseline model
- Calculates the distance between similar data points, and makes predictions based on the frequency of different classes



# KNN Baseline Model

- Build a baseline KNN model with hyperparameters
  - `n_neighbors = 5`, `p = 2`, `leaf_size = 30`
- Resulted in an initial accuracy of 61% on the training data

	precision	recall	f1-score	support
0.0	0.74	0.74	0.74	82208
1.0	0.24	0.24	0.24	28421
accuracy			0.61	110629
macro avg	0.49	0.49	0.49	110629
weighted avg	0.61	0.61	0.61	110629

# KNN Hyperparameter Tuning

- Use Grid Search CV to tune the hyperparameters
- Find the hyperparameters for the highest accuracy model

```
# Build a dictionary for the hyperparameters we are looking to tune
# ignore p value and leaf_size as gridsearchcv takes too long, so just tune n_neighbors
# p = [1, 2]
#leaf_size = list(range(1, 50))
n_neighbors = list(range(1, 50))
hyperparameters = dict(n_neighbors = n_neighbors)
```

```
# initiate another KNN model class
knn_2 = KNeighborsClassifier()

#run gridsearchcv on the hyperparameters for a KNN model with 10 folds
clf = GridSearchCV(knn_2, hyperparameters, cv = 10, n_jobs = -1)

# fit the model
best_model = clf.fit(X_train_pca, Y_train)
```

```
print('Best leafsize:' , best_model.best_estimator_.get_params()['leaf_size'])
print('Best n_neighbors:' , best_model.best_estimator_.get_params()['n_neighbors'])
#print('Best p:' , best_model.best_estimator_.get_params()['p'])
```

```
Best leafsize: 30
Best n_neighbors: 27
```

# KNN Model Results

- Using the tuned hyperparameters, a KNN model fit on training data results in an accuracy on the test data of 67%
- Precision, recall, and F1-scores are 74%, 84%, and 79% respectively

	precision	recall	f1-score	support
0.0	0.74	0.84	0.79	82208
1.0	0.27	0.17	0.21	28421
accuracy			0.67	110629
macro avg	0.51	0.50	0.50	110629
weighted avg	0.62	0.67	0.64	110629

# Model - Random Forest

## Building a random forest

```
1 %%time
2 random_forest = RandomForestClassifier(criterion='entropy', #function to measure the quality of a feature split
3                                     n_estimators=10, #number of trees in the forest
4                                     bootstrap = True, #different sample for each estimator
5                                     max_samples = 0.6, #max number of samples is 50% of training data
6                                     n_jobs=2) # The number of jobs to run in parallel.
7 clf = random_forest.fit(X_train, Y_train)
```

Accuracy on Training Data: 0.9724  
Accuracy on Validation Data: 0.8727  
Validation Precision score: 0.8360  
Validation Recall score: 0.8239  
Validation F1 score 0.8296

	0	1
0	2488903	20304
1	71148	738515
Predicted label		

# Random Forest - Hyperparameter Tuning

- Establish random forest parameter distribution
  - SKlearn's RandomizedSearchCV will randomly sample from these

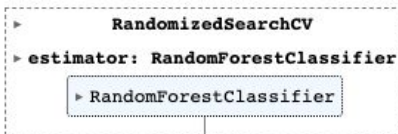
```
1 criterion = ['entropy'] #what will be used to choose features at node splits
2 n_estimators = [i for i in range(10,18,2)] # number of trees in the random forest
3 bootstrap = [True] # drawing samples from our source data with replacement
4 max_samples = [i/10 for i in range(2,10,2)] #percentage of training samples to train each tree with
5 n_jobs=[-1]
6 max_features = ['sqrt'] #could add in 'log2' but took too long
7 max_depth = [i for i in range(110,160,10)] #max amount of layers in a tree
8 max_depth.append(None) #if None, then nodes are expanded until all leaves are pure
9 min_samples_split = [2, 6, 10] # minimum sample number to split a node, need at least two to split
10 min_samples_leaf = [1,2, 3] # A split point at any depth will only be considered if it leaves at least
11                               #min_samples_leaf training samples in each of the left and right branches.
12                               #This may have the effect of smoothing the model, especially in regression
13
14 random_grid = { 'criterion' : criterion,
15                 'n_estimators' : n_estimators,
16                 'bootstrap' : bootstrap,
17                 'max_samples' : max_samples,
18                 'n_jobs' : n_jobs,
19                 'max_features' : max_features,
20                 'max_depth' : max_depth,
21                 'min_samples_split': min_samples_split,
22                 'min_samples_leaf' : min_samples_leaf
23 }
```

# Random Forest - Hyperparameter Tuning

- Initiate the tuning

```
1 %%time
2 #first we need to initiate a new base tree estimator
3 base_tree = RandomForestClassifier()
4
5 #initialize the random search CV hyperparameter tuner
6 Random_Forest_HyperTuned = RandomizedSearchCV(estimator = base_tree,
7         param_distributions = random_grid,#params defines above
8         n_iter = 8, #Number of parameter settings that are sampled from distri
9         cv = 3, #default is 5-fold cross validation, so we will have 30 trees
10        verbose = 1, #the higher, the more messages
11        return_train_score = True # attribute will include training scores
12        )
13 #now we need to train the model using our training data and labels
14 Random_Forest_HyperTuned.fit(X_train,Y_train)
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits  
CPU times: user 21min 23s, sys: 2min 17s, total: 23min 40s  
Wall time: 12h 50min 54s



## Best parameters

```
{'n_jobs': -1,  
 'n_estimators': 16,  
 'min_samples_split': 2,  
 'min_samples_leaf': 3,  
 'max_samples': 0.8,  
 'max_features': 'sqrt',  
 'max_depth': 150,  
 'criterion': 'entropy',  
 'bootstrap': True}
```

## Feature Importances

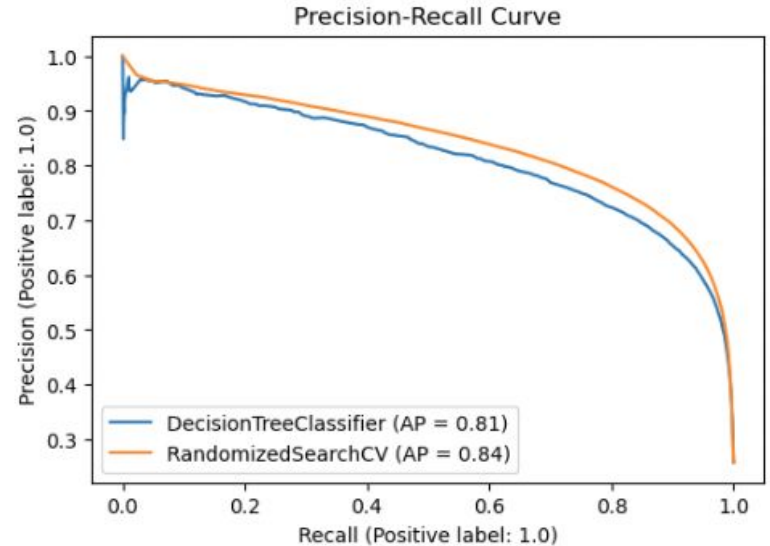
	Feature	importance
0	P_2	0.076980
1	D_44	0.038740
2	D_48	0.032910
3	D_45	0.029680
4	D_61	0.027010
5	B_20	0.025950
6	B_3	0.021830
7	B_2	0.020580
8	B_10	0.019860
9	B_18	0.018480



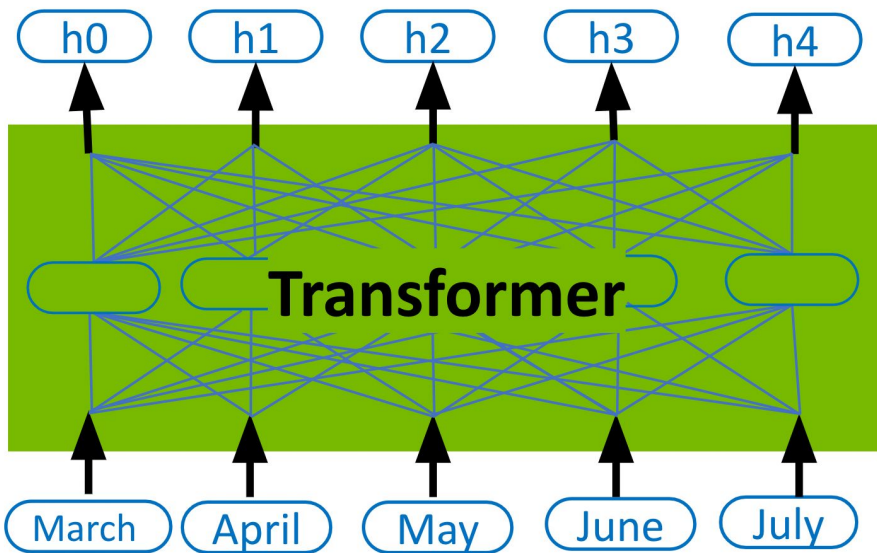
# Random Forest - Test Set Evaluation

Accuracy on Test Data: 0.8829  
Test Precision score: 0.8422  
Test Recall score: 0.8641  
Test F1 score 0.8520

True label	0	1
0	2488903	20304
1	71148	738515
Predicted label		



# Model Transformer

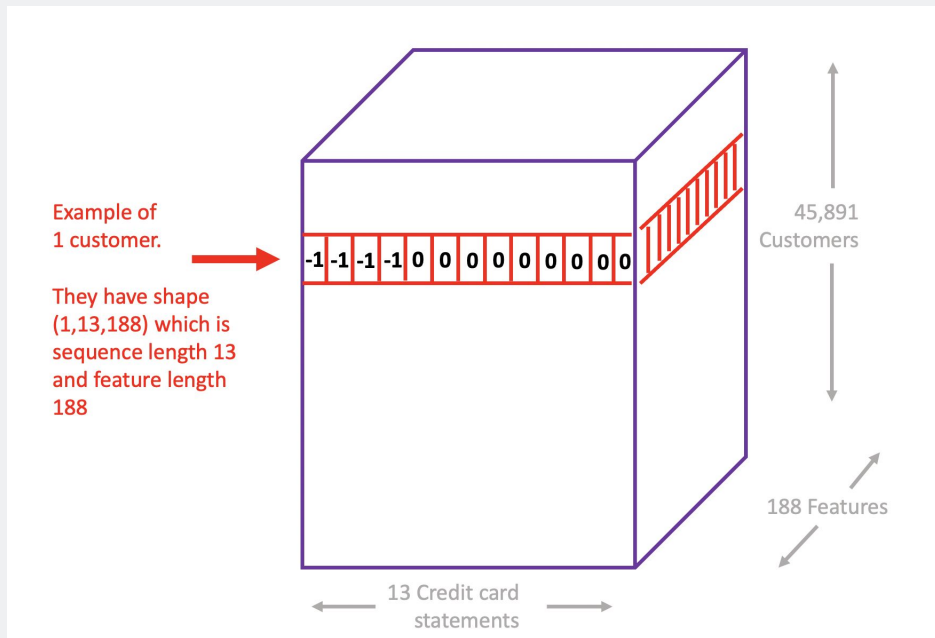


Time series data:

- Most of customers have 13 months credit history
- Transform data from 2d to 3d
- Implement transformer encoder and connect its output to a classification head using binary cross-entropy

# Model Transformer - Pre-Processing and Feature Engineering

- Load and split training data:
  - 10 files
  - 45891 customers
  - 13 credit card statements
- Add padding
  - For customers has less than 13 credit card history, pad with -1
- Feature Engineer
  - Output shape (45891, 13, 188)



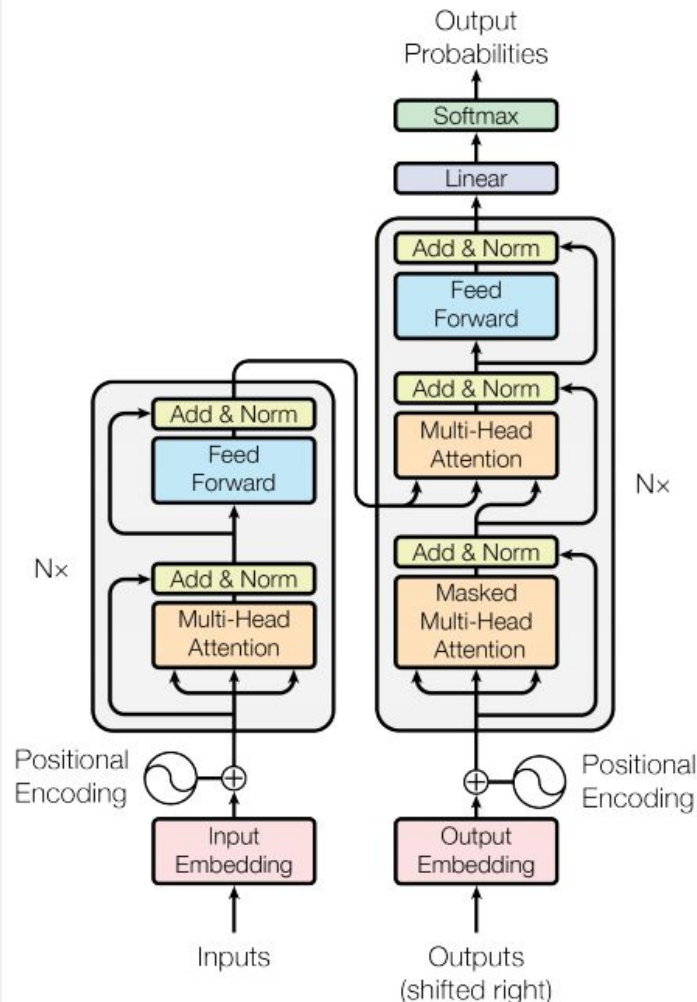
# Transformer Block

## Hyper-parameters

```
feat_dim = 188
embed_dim = 64 # Embedding size for attention
num_heads = 4 # Number of attention heads
ff_dim = 128 # Hidden layer size in feed forward network inside transformer
dropout_rate = 0.3
num_blocks = 2
```

```
[ ] class TransformerBlock(layers.Layer):
    def __init__(self, embed_dim, feat_dim, num_heads, ff_dim, rate=0.1):
        super(TransformerBlock, self).__init__()
        self.att = layers.MultiHeadAttention(num_heads=num_heads, key_dim=embed_dim)
        self.ffn = keras.Sequential(
            [layers.Dense(ff_dim, activation="gelu"), layers.Dense(feat_dim),]
        )
        self.layernorm1 = layers.LayerNormalization(epsilon=1e-6)
        self.layernorm2 = layers.LayerNormalization(epsilon=1e-6)
        self.dropout1 = layers.Dropout(rate)
        self.dropout2 = layers.Dropout(rate)

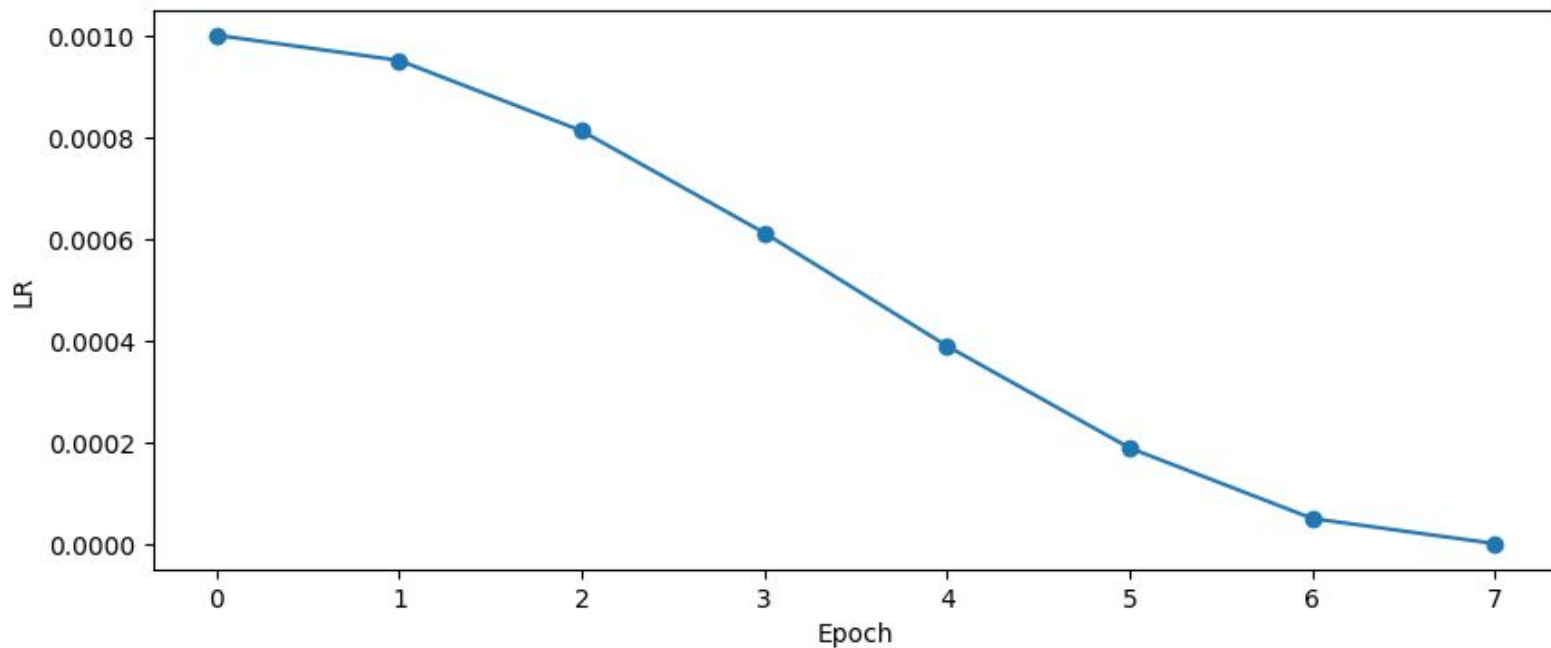
    def call(self, inputs, training):
        attn_output = self.att(inputs, inputs)
        attn_output = self.dropout1(attn_output, training=training)
        out1 = self.layernorm1(inputs + attn_output)
        ffn_output = self.ffn(out1)
        ffn_output = self.dropout2(ffn_output, training=training)
        return self.layernorm2(out1 + ffn_output)
```



# Transformer (cont.)

- Learning schedule

↳ Learning rate schedule: 0.001 to 0.001 to 1e-06



# Transformer - Training

- We train 5 folds of model using cross validation
  - Use 8 files to train, 2 files for validation
- 8 epochs for each fold
- Save 5 output transformer model into files

```
#####
```

```
### Fold 5 with valid files [9, 10]
```

```
### Training data shapes (367128, 13, 188) (367128,)
```

```
### Validation data shapes (91785, 13, 188) (91785,)
```

```
#####
```

```
Epoch 1: LearningRateScheduler setting learning rate to 0.001.
```

```
Epoch 1/8
```

```
718/718 - 24s - loss: 0.2424 - val_loss: 0.2307 - lr: 0.0010 - 24s/epoch - 34ms/step
```

```
Epoch 2: LearningRateScheduler setting learning rate to 0.0009505339495172585.
```

```
Epoch 2/8
```

```
718/718 - 19s - loss: 0.2312 - val_loss: 0.2370 - lr: 9.5053e-04 - 19s/epoch - 27ms/step
```

```
Epoch 3: LearningRateScheduler setting learning rate to 0.0008119331560284375.
```

```
Epoch 3/8
```

```
718/718 - 19s - loss: 0.2279 - val_loss: 0.2269 - lr: 8.1193e-04 - 19s/epoch - 27ms/step
```

# Model Transformer - Performance

- Kaggle Evaluation Metrics
  - The metics M is defined as the the mean of two measures
  - G: Normalized Gini Coefficient
  - D: Default rate captured at 4%

$$M = 0.5 \cdot (G + D)$$

## Cross Validation Score

Fold 1 CV= 0.7843216315379575

Fold 2 CV= 0.7821227060205045

Fold 3 CV= 0.7852668054694807

Fold 4 CV= 0.7878325830028957

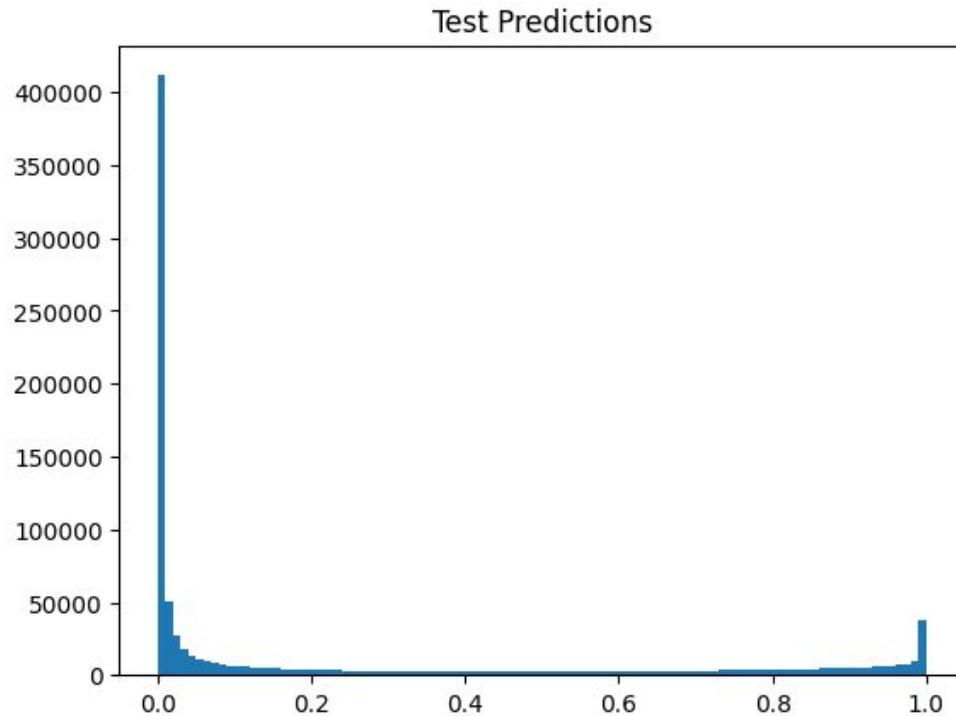
Fold 5 CV= 0.7903108125684797

#####

Overall CV = 0.7859054438364497

# Model Transformer - Test Inference

- Process test data the same way as train data, the output will be 20 files with each file having shape (46231, 13, 188) as (customer x statement x feature)
- Ensemble 5 models from training and create predictions by averaging each model outputs
- Graph the predictions





# Model Transformer - Kaggle submission

- Private score: 0.798, Public Score: 0.789

Submission and Description

Private Score ⓘ

Public Score ⓘ

Selected



submission.csv





















0.79831

0.78911



Complete (after deadline) · 3d ago · Transformer model

■ Prize Winners

#	△	Team	Members	Score	Entries	Last	Solution
1	▲ 3	lucky shake		0.80977	91	8mo	
2	▲ 5	byDefault [JuneHomes]	   	0.80938	340	8mo	
3	▲ 2	AIBANK		0.80925	50	8mo	
4	▲ 6	GRE	    	0.80900	186	8mo	
5	▼ 3		   	0.80881	387	8mo	

## Takeaways

- No PII data in features
  - No insight into categorical features
- Models may be learning abstract biases about certain demographics
- Last 13 bank statements may introduce recency bias

## Contributions:

- EDA (**All**)
- Data Pre Processing (**All**)
- KNN (**Jim**)
- Random Forest (**Julian**)
- Transformer (**Rick**)